


SYSTEM CALLS

Jevitha K.P

Secure Coding Lab 3



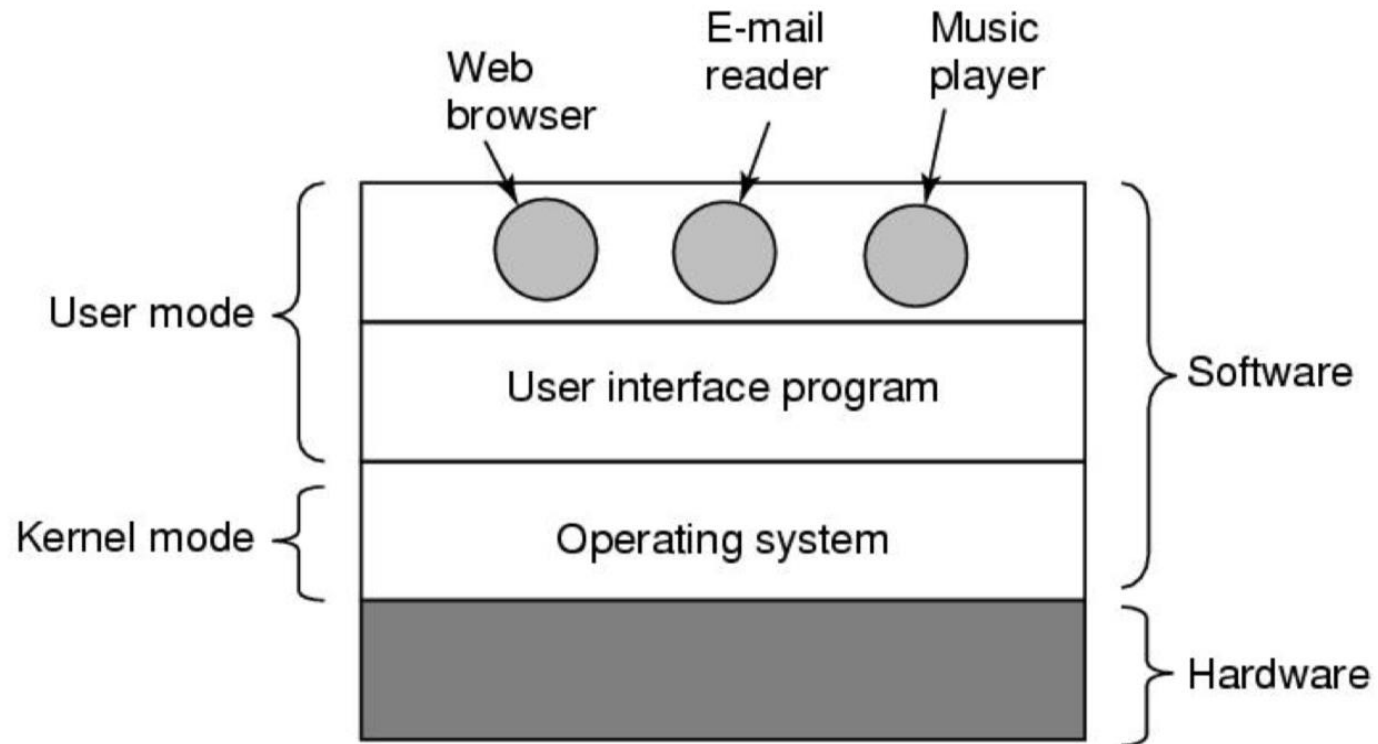
SYSTEM CALLS

- System Calls is the mechanism used by an application program to **request services from the operating system (kernel)** such as opening a file, forking to a new process, or requesting more memory.
- It is the **real process to kernel communication** mechanism used by all processes
- System calls often use a **special machine code instruction** which causes the processor to change mode (e.g. to "supervisor mode" or "protected mode").
- This allows the OS to perform restricted actions such as accessing hardware devices or the memory management unit.

SYSTEM CALLS

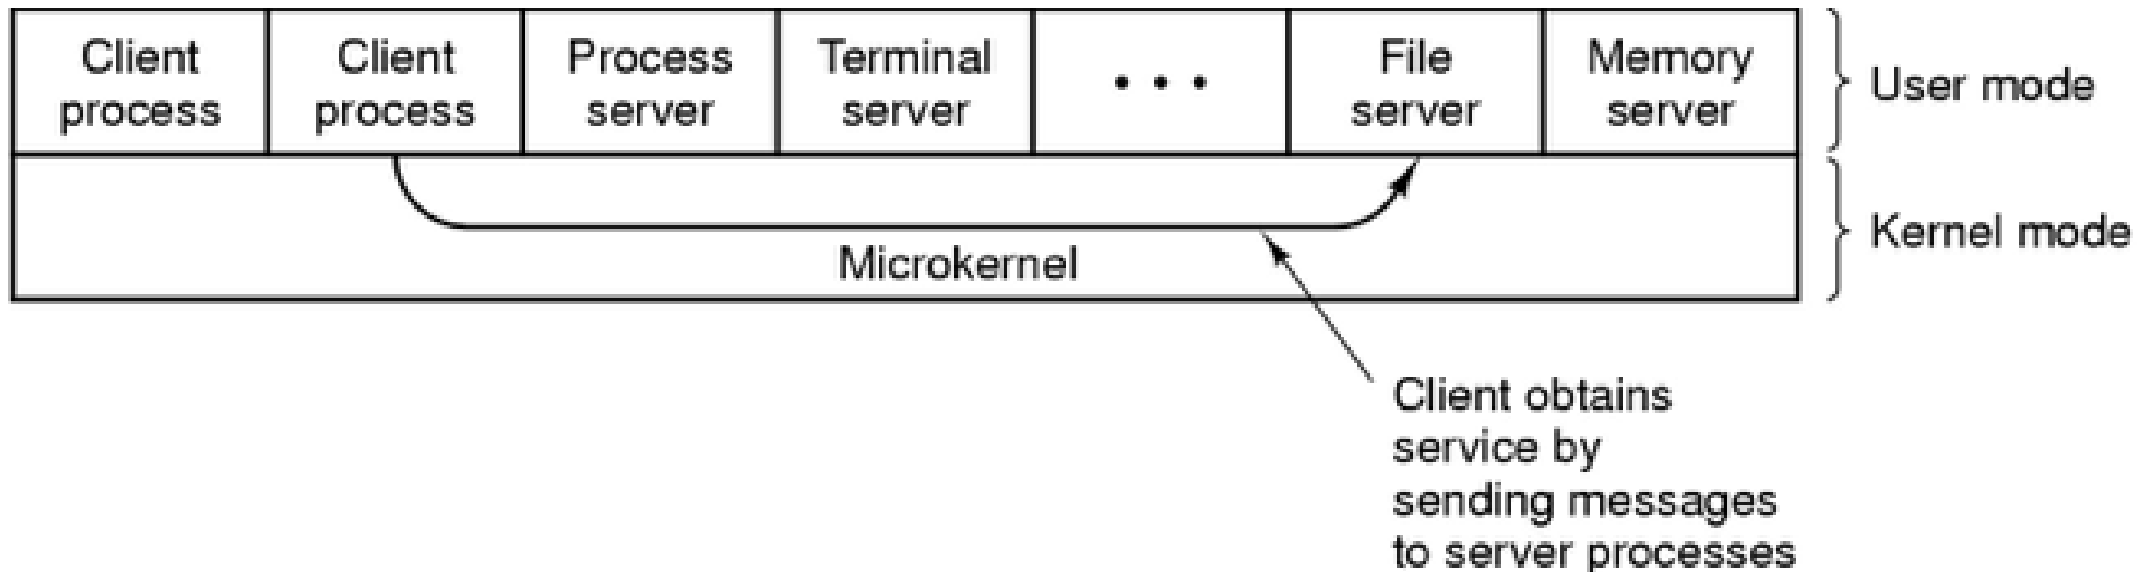
- A process, in general, is not supposed to be able to access the kernel.
- It can't access kernel memory and can't call kernel functions , which is enforced by the hardware of the CPU (hence, called '**protected mode**').
- System calls are an exception to this general rule.

USER MODE VS KERNEL MODE



CLIENT-SERVER MODEL INSIDE OS

- The OS may be split into
 - a **kernel** which is always present
 - various **system programs** use facilities provided by the kernel to perform higher-level house-keeping tasks,
 - often acting as servers in a client-server relationship.



SYSTEM CALLS IN LINUX

- In Linux, System call service is provided by Linux kernel.
- In C programming, functions defined in **libc library** provides a wrapper for many system calls.
- It is also possible to invoke **syscall()** function directly.
- Each system call has a function number defined in **<syscall.h>** or **<unistd.h>**.
- Internally, system call is invoked by software interrupt **0x80** to transfer control to the kernel.
- System call table is defined in Linux kernel source file **“arch/i386/kernel/entry.S ”**.


STEPS IN MAKING A SYSTEM CALL

- There are several steps in making the system call :
 - The calling process fills the registers with the appropriate values
 - It then sends software **interrupt 0x80**, which jumps to a previously defined location in the kernel (the location is readable by user processes, but not writable by them).
- The hardware knows that once we jump to this location,
 - we are no longer running in restricted user mode,
 - but as the operating system kernel --- and therefore allowed to do the task we want.

STEPS IN MAKING A SYSTEM CALL

read (fd, buffer, nbytes)

- 1-3 - push 3 parameters on stack
- 4 - invoke the system call
- 5 - put code for system call on register
- 6 - trap to the kernel
- 7-10 Since a number is associated with each system call, **system call interface** invokes/dispatch intended system call in OS kernel and return status of the system call and any return value
- 11 - increment stack pointer



SYSTEM CALLS

- The location in the kernel a process can jump to is called *system_call*.
- The procedure at that location **checks the system call number**, which tells the kernel what service the process requested.
- Then, it looks at the **table of system calls (sys_call_table)** to see the address of the kernel function to call.
- Then it calls the function, and after it returns, does a few system checks and then returns back to the process (or to a different process, if the process time ran out).

SYSTEM CALL ERROR

- When a system call results in an error, it returns -1 and stores the reason the call failed in an external variable named "**errno**".
- When a system call returns **successfully**, it returns **something other than -1**, but it does not clear "errno".
- "errno" only has meaning directly after a system call that returns an error.
- The "**/usr/include/errno.h**" file maps these error numbers to manifest constants, and it is these constants that we should use in our programs.

SYSTEM CALL CATEGORIES

- **System calls are divided into 5 categories mainly :**
- Process Management
- File / Directory Management
- Device Management
- Information Maintenance
- Communication

PROCESS MANAGEMENT

- This system calls perform the task of process creation, process termination, etc.
- The Linux System calls under this are **fork()** , **exit()** , **exec()**.

Process management

Call	Description
<code>pid = fork()</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

FILE MANAGEMENT

- File management system calls handle file manipulation jobs like creating a file, reading, and writing, etc.
- Some calls under this are **open()**, **read()**, **write()**, **close()**.

File management	
Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &buf)</code>	Get a file's status information

DIRECTORY MANAGEMENT

Directory and file system management

Call	Description
<code>s = mkdir(name, mode)</code>	Create a new directory
<code>s = rmdir(name)</code>	Remove an empty directory
<code>s = link(name1, name2)</code>	Create a new entry, name2, pointing to name1
<code>s = unlink(name)</code>	Remove a directory entry
<code>s = mount(special, name, flag)</code>	Mount a file system
<code>s = umount(special)</code>	Unmount a file system

DEVICE MANAGEMENT

- Involves tasks like reading from device buffers, writing into device buffers, etc.
- The Linux System calls under this are `ioctl()`, `read()`, `write()`.
- **`ioctl()`:**
 - `ioctl()` is referred to as Input and Output Control.
 - `ioctl` is a system call for device-specific input/output operations and other operations which cannot be expressed by regular system calls.

INFORMATION MAINTENANCE

- It handles information and its transfer between the OS and the user program.
- In addition, OS keeps the information about all its processes and system calls are used to access this information.
- The System calls under this are **getpid()**, **alarm()**, **sleep()**.
- **getpid()**
 - getpid stands for Get the Process ID.
 - The getpid() function shall return the process ID of the calling process.
 - The getpid() function shall always be successful and no return value is reserved to indicate an error.

INFORMATION MAINTENANCE

- **alarm()**

- This system call sets an alarm clock for the delivery of a signal that when it has to be reached.
- It arranges for a signal to be delivered to the calling process.

- **sleep()**

- This System call suspends the execution of the currently running process for some interval of time
- Meanwhile, during this interval, another process is given chance to execute

COMMUNICATION

- These types of system calls are specially used for inter-process communications.
- Two models are used for inter-process communication
 - Message Passing(processes exchange messages with one another)
 - Shared memory(processes share memory region to communicate)
- The system calls under this are **pipe()** , **shmget()** ,**mmap()**.
- **pipe():**
 - The pipe() system call is used to communicate between different Linux processes.
 - It is mainly used for inter-process communication.
 - The pipe() system function is used to open file descriptors.

COMMUNICATION

- **shmget():**

- shmget stands for shared memory segment.
- It is mainly used for Shared memory communication.
- This system call is used to access the shared memory and access the messages in order to communicate with the process.

- **mmap():**

- This function call is used to map or unmap files or devices into memory.
- The mmap() system call is responsible for mapping the content of the file to the virtual memory space of the process.

SYSTEM CALLS FOR MISCELLANEOUS TASKS

Miscellaneous

Call	Description
<code>s = chdir(dirname)</code>	Change the working directory
<code>s = chmod(name, mode)</code>	Change a file's protection bits
<code>s = kill(pid, signal)</code>	Send a signal to a process
<code>seconds = time(&seconds)</code>	Get the elapsed time since Jan. 1, 1970

FILE DESCRIPTORS

- Each UNIX process has 20 file descriptors at its disposal, numbered 0 through 19.
- The first three are already opened when the process begins
 - 0: The standard input
 - 1: The standard output
 - 2: The standard error output
- When the parent process forks a process, the child process inherits the file descriptors of the parent.

LET'S MOVE TO PROGRAMMING ...

SYSCALL, GETPID

```
#include <syscall.h>
#include <unistd.h>
#include <stdio.h>
#include <sys/types.h>
int main(void) {
    long ID1, ID2;
    /* direct system call - SYS_getpid (func
    no. is 20) */
    ID1 = syscall(SYS_getpid);
    printf ("syscall(SYS_getpid)=%ld\n",
    ID1);
```

```
/* "libc" wrapped system call */
/* SYS_getpid (Func No. is 20) */
ID2 = getpid();
printf ("getpid()=%ld\n", ID2);
return(0);
}
```

FORK

- `int fork()` turns a single process into 2 identical processes, known as the **parent** and the **child**.
- On success, `fork()` returns 0 to the child process and returns the process ID of the child process to the parent process.
- The child process will have its own unique PID.
- On failure, `fork()` returns -1 to the parent process, sets `errno` to indicate the error, and no child process is created.

FORK

- Simple use of fork, where two copies are made and run together (multitasking)

-

- `main()`

- `{`

```
int return_value;
```

```
printf("Forking process\n");
```

```
return_value = fork();
```

```
printf("The process id is %d and return value is %d\n", getpid(), return_value);
```

```
    printf("This line is not printed\n");
```

```
} Secure Coding Lab
```

EXEC

Exec.c

```
▪ #include <stdio.h>
▪ #include <unistd.h>
▪ #include <stdlib.h>
▪ int main(int argc, char *argv[])
▪ {
▪     printf("PID = %d\n", getpid());
▪     char *args[] = {"Hello", "C", "Programming", NULL};
▪     execv("./hello", args);
▪     printf("Back to exec.c - This line will not be executed");
▪     return 0;
▪ }
```

Hello.c

```
▪ #include <stdio.h>
▪ #include <unistd.h>
▪ #include <stdlib.h>
▪ int main(int argc, char *argv[])
▪ {
▪     printf("We are in Hello.c\n");
▪     printf("PID of hello.c = %d\n", getpid());
▪     return 0;
▪ }
> gcc hello.c -o hello
```

STRACE

- If we want to see which system calls a program uses, the **strace** command is used :
- **strace <arguments>.**
- Execute the programs we did using strace.